

# Yelp Dataset Experiments

February 11, 2021

## 1 Introduction

This document includes analysis and experiments carried out on [Yelp Dataset](#). It contains reviews and corresponding users and businesses in North America. We first do some exploratory data analysis in the Section 2. Then we do a few machine learning tasks explained in Section 3. In this report, most of the code cells and outputs are removed for clarity. If you're interested, you could check the accompanying jupyter-notebook files.

## 2 Data Analysis

We first analyze the businesses.

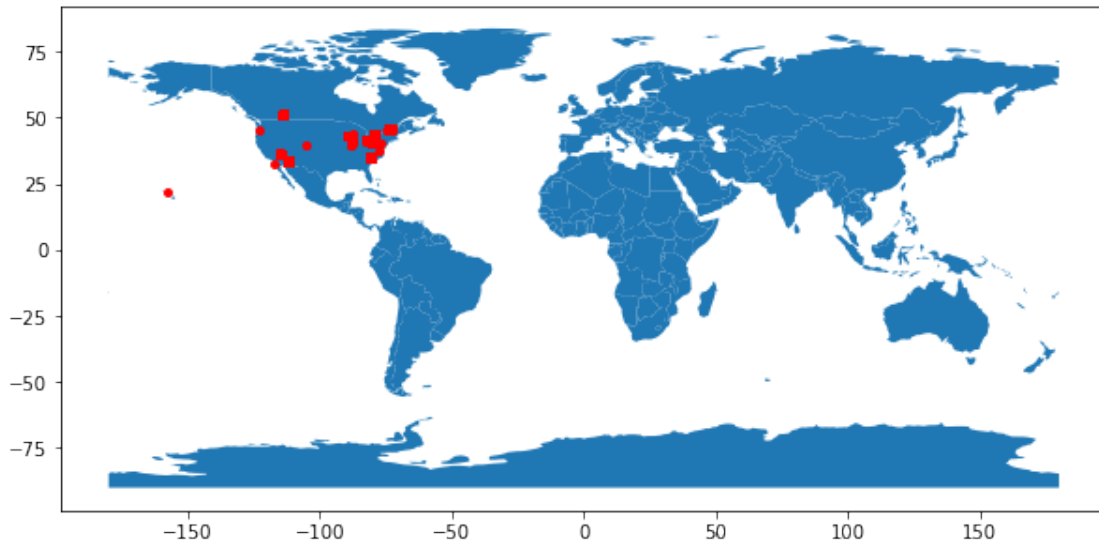
Let's import necessary modules and read the business file.

```
[9]: import pandas as pd
import matplotlib.pyplot as plt
from shapely.geometry import Point
import geopandas as gpd
from geopandas import GeoDataFrame
```

Using latitude and longitude, we plot the map and see the locations of business.

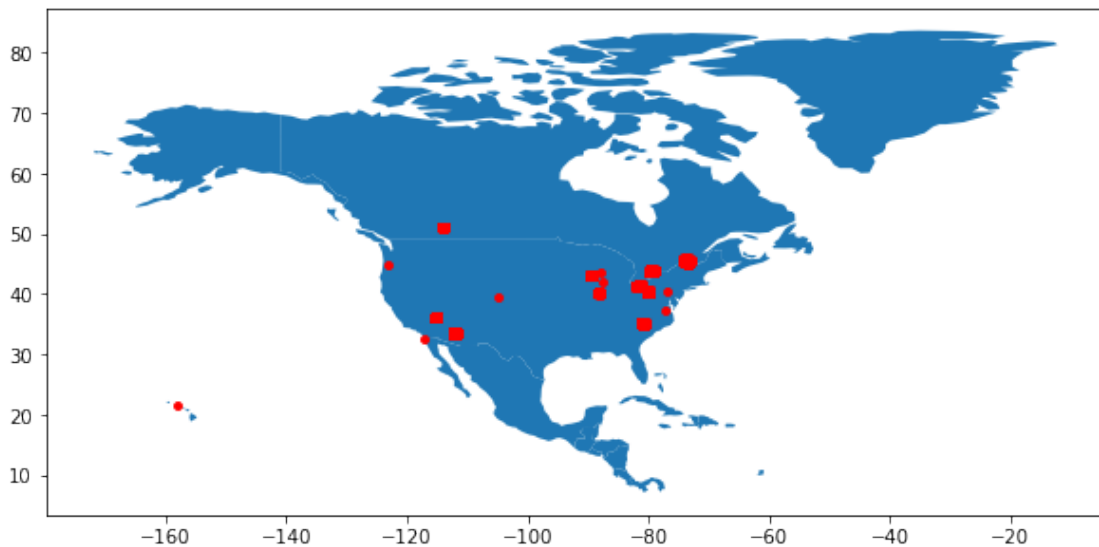
```
[11]: geometry = [Point(xy) for xy in zip(df['longitude'], df['latitude'])]
gdf = GeoDataFrame(df, geometry=geometry)

#this is a simple map that goes with geopandas
world = gpd.read_file(gpd.datasets.get_path('naturalearth_lowres'))
gdf.plot(ax=world.plot(figsize=(10, 6)), marker='o', color='red', markersize=15);
```



```
[12]: world = world[world.continent == 'North America']
      gdf.plot(ax=world.plot(figsize=(10, 6)), marker='o', color='red', markersize=15);
```

They are mainly located in North America, so we zoom a little



```
[13]: world = world[world.name == 'United States of America']
      gdf.plot(ax=world.plot(figsize=(10, 6)), marker='o', color='red', markersize=15);
```

They are mainly located in United States, so we zoom a little more

There are 209393 rows (businesses). Attributes, categories and hours are not provided for some

**businesses**

```
[16]: null_rows = df[df.isnull().any(axis=1)]
      select_columns = ['business_id', 'name', 'stars', 'review_count', 'is_open',
                        'attributes', 'categories', 'hours']
      null_rows[select_columns].head(10)
```

Attributes, categories and hours are strings and null values cannot directly be filled with average statistics from other rows.

```
[17]: # check examples of other columns
      list_of_columns = df.columns.tolist()
      print(list_of_columns)
```

```
['business_id', 'name', 'address', 'city', 'state', 'postal_code', 'latitude',
'longitude', 'stars', 'review_count', 'is_open', 'attributes', 'categories',
'hours', 'geometry']
```

```
[18]: df[list_of_columns[:-3]].head()
```

The majority of business are located in Las Vegas, Toronto, Phoenix, Charlotte, Scottsdale and cities like Queensville, Huntigdon, Rocky River, Rainbow Valley and Gilbert, among others have listed only 1 business each.

```
[20]: # stats for number of reviews
      df.review_count.describe()
```

```
[20]: count      209393.000000
      mean         36.937505
      std         123.343597
      min           3.000000
      25%           4.000000
      50%           9.000000
      75%          27.000000
      max        10129.000000
      Name: review_count, dtype: float64
```

The minimum number of reviews for a business listed in the dataset is 3. 75% of the the businesses have received 27 reviews or less. That means very few businesses have received extremely high number of reviews thus skewing the mean, i.e., 36 reviews. First we check how many businesses have fewer than 50 reviews out of the 209393 business.

```
[21]: df[df['review_count']<50].shape[0]
```

```
[21]: 177451
```

177451 out of 209393 business have fewer than 50 reviews leaving 31942 with higher number of reviews than 50.

Now, let's 10 businesses with the least number of reviews.

```
[22]: df[['name', 'city', 'state', 'review_count', 'stars']].
      ↪sort_values(by='review_count').head(10)
```

```
[22]:
```

	name	city	state	review_count	\
209392	Kudlow Ye	Toronto	ON	3	
190418	241 Pizza	Scarborough	ON	3	
75451	MAI Montreal Arts Interculturels	Montréal	QC	3	
75454	Sachika	Montréal	QC	3	
75461	Jungle Juice	Toronto	ON	3	
190414	SAS Too	Phoenix	AZ	3	
75467	Toepel Company	Mesa	AZ	3	
75468	Baja Ready Mix	Phoenix	AZ	3	
75471	F T Financial	Scottsdale	AZ	3	
75476	Grand and Clover Cocktail Co.	Toronto	ON	3	

	stars
209392	5.0
190418	3.5
75451	4.5
75454	4.5
75461	3.5
190414	2.0
75467	5.0
75468	3.5
75471	5.0
75476	5.0

And 10 businesses with the most number of reviews.

```
[23]: df[['name', 'city', 'state', 'review_count', 'stars']].
      ↪sort_values(by='review_count', ascending=False).head(10)
```

```
[23]:
```

	name	city	state	review_count	stars
81545	Bacchanal Buffet	Las Vegas	NV	10129	4.0
118008	Mon Ami Gabi	Las Vegas	NV	9264	4.0
147379	Wicked Spoon	Las Vegas	NV	7383	3.5
83020	Hash House A Go Go	Las Vegas	NV	6751	4.0
201975	Gordon Ramsay BurGR	Las Vegas	NV	5494	4.0
95962	Earl of Sandwich	Las Vegas	NV	5232	4.5
22754	Yardbird Southern Table & Bar	Las Vegas	NV	4828	4.5
145294	Secret Pizza	Las Vegas	NV	4803	4.0
205740	The Buffet At Wynn	Las Vegas	NV	4803	3.5
77432	The Cosmopolitan of Las Vegas	Las Vegas	NV	4740	4.0

Let's aggregate the number of reviews for each city.

We first list 20 cities with the smallest total reviews for the businesses located there.

```
[24]: grouped_review = df.groupby(['city']).agg({'review_count': ['sum', 'mean', 'min', 'max']})
grouped_review.columns = ['review_sum', 'review_mean', 'review_min', 'review_max']
grouped_review = grouped_review.reset_index()
grouped_review.sort_values(by='review_sum').head(20)
```

```
[24]:
```

	city	review_sum	review_mean	review_min	\
553	McMasterville	3	3.0	3	
1048	St-Laurent	3	3.0	3	
1045	St-Jean Sur Richelieu	3	3.0	3	
1044	St-Clet	3	3.0	3	
316	Fountain Hls	3	3.0	3	
167	Chateau	3	3.0	3	
1029	South Heights	3	3.0	3	
418	Joliet	3	3.0	3	
1025	Somerton	3	3.0	3	
1049	St-Lazare	3	3.0	3	
735	O'hara Township	3	3.0	3	
176	Citibank	3	3.0	3	
1014	Sharon	3	3.0	3	
179	Claremont	3	3.0	3	
1007	Sgs Industrial Park	3	3.0	3	
183	Cleveland	3	3.0	3	
1002	Scottsdsale	3	3.0	3	
742	Oakvile	3	3.0	3	
408	Indian Land, South Carolina	3	3.0	3	
174	Chomedey, Laval	3	3.0	3	

```
review_max
553      3
1048     3
1045     3
1044     3
316      3
167      3
1029     3
418      3
1025     3
1049     3
735      3
176      3
1014     3
179      3
1007     3
183      3
1002     3
```

742	3
408	3
174	3

We then list 20 cities with the highest total reviews for the businesses located there.

```
[25]: grouped_review.sort_values(by='review_sum').tail(20)
```

```
[25]:
```

	city	review_sum	review_mean	review_min	review_max
521	Markham	59004	30.058074	3	638
705	North Las Vegas	60405	36.019678	3	815
599	Mississauga	65557	18.634736	3	1009
792	Peoria	69539	33.432212	3	713
135	Calgary	107106	12.785723	3	512
508	Madison	114475	31.065129	3	1879
341	Glendale	118096	30.882845	3	1075
182	Cleveland	129489	33.572466	3	1372
337	Gilbert	153826	41.075033	3	2369
626	Montréal	171059	24.510532	3	2667
158	Chandler	183475	40.060044	3	1346
577	Mesa	197214	29.985404	3	1362
1095	Tempe	229594	47.861997	3	2400
376	Henderson	261244	49.553111	3	2456
826	Pittsburgh	262412	34.392136	3	2001
163	Charlotte	371580	35.653425	3	2174
1000	Scottsdale	439439	47.039071	3	2369
1106	Toronto	583512	28.651282	3	2758
804	Phoenix	842321	41.759010	3	3515
461	Las Vegas	2360735	74.633587	3	10129

Las Vegas, Toronto, Phoenix, Charlotte, and Scottsdale has higher total reviews than other cities which is understandable since there are more businesses in the dataset located in these locations, as shown above. Let's show a scatter plot of number of businesses and number of reviews for each city.

We first merged the two dataframes for number of total reviews and number of businesses.

```
[26]: grouped_business = df.value_counts('city').reset_index()
grouped_business.columns= ['city', 'num_businesses']
grouped_business
```

```
[26]:
```

	city	num_businesses
0	Las Vegas	31631
1	Toronto	20366
2	Phoenix	20171
3	Charlotte	10422
4	Scottsdale	9342
...	...	...

1246	Queensville	1
1247	Huntingdon	1
1248	ROCKY RIVER	1
1249	Rainbow Valley	1
1250	Gilbert	1

[1251 rows x 2 columns]

```
[27]: merged = pd.merge(grouped_review, grouped_business, on='city')
merged
```

```
[27]:
```

	city	review_sum	review_mean	review_min	review_max	\
0		7	3.5	3	4	
1	110 Las Vegas	43	43.0	43	43	
2	4321 W Flamingo Rd	292	292.0	292	292	
3	ARSENAL	4	4.0	4	4	
4	AZ	4	4.0	4	4	
...	...	...	...	...	...	
1246	toronto	24	8.0	3	11	
1247	Île-Perrot	8	8.0	8	8	
1248	Île-des-Soeurs	5	5.0	5	5	
1249	Avondale	5	5.0	5	5	
1250	Gilbert	17	17.0	17	17	


	num_businesses
0	2
1	1
2	1
3	1
4	1
...	...
1246	3
1247	1
1248	1
1249	1
1250	1

[1251 rows x 6 columns]

**We can now show the scatter plot.**

```
[28]: merged.plot(kind='scatter', x='review_sum', y='num_businesses')
```

```
[28]: <AxesSubplot:xlabel='review_sum', ylabel='num_businesses'>
```




Taking a first look at the scatterplot, it seems there is a positive correlation between the number of reviews and the number of businesses. There are some extreme outlier cities with 2000-3000 businesses. Let's just plot only those with fewer than 9000 businesses in total.

```
[29]: merged[merged['num_businesses']<9000].plot(kind='scatter', x='review_sum',  
        ↳y='num_businesses')
```

```
[29]: <AxesSubplot:xlabel='review_sum', ylabel='num_businesses'>
```






The trend is still the same, and it makes sense. The more businesses a city has, the more people it contains and thus, the higher total reviews for businesses in that city. It is not that people in some cities are less engaging.

Let's also check the star ratings given to the businesses.

```
[30]: df['stars'].plot(kind='hist', bins=10)
```


```
[30]: <AxesSubplot:ylabel='Frequency'>
```



In general, businesses have received 4 star ratings. Let's see how that correlates to the number of reviews.


```
[31]: df.plot(kind='scatter', x='review_count', y='stars')
```

```
[31]: <AxesSubplot:xlabel='review_count', ylabel='stars'>
```



```
[32]: df.plot(kind='scatter', x='stars', y='review_count')
```

```
[32]: <AxesSubplot: xlabel='stars', ylabel='review_count'>
```



In general, it does not depend on how many reviews a business received. The star rating could actually reflect the general public liking of the place.

```
[34]: df['categories']
```

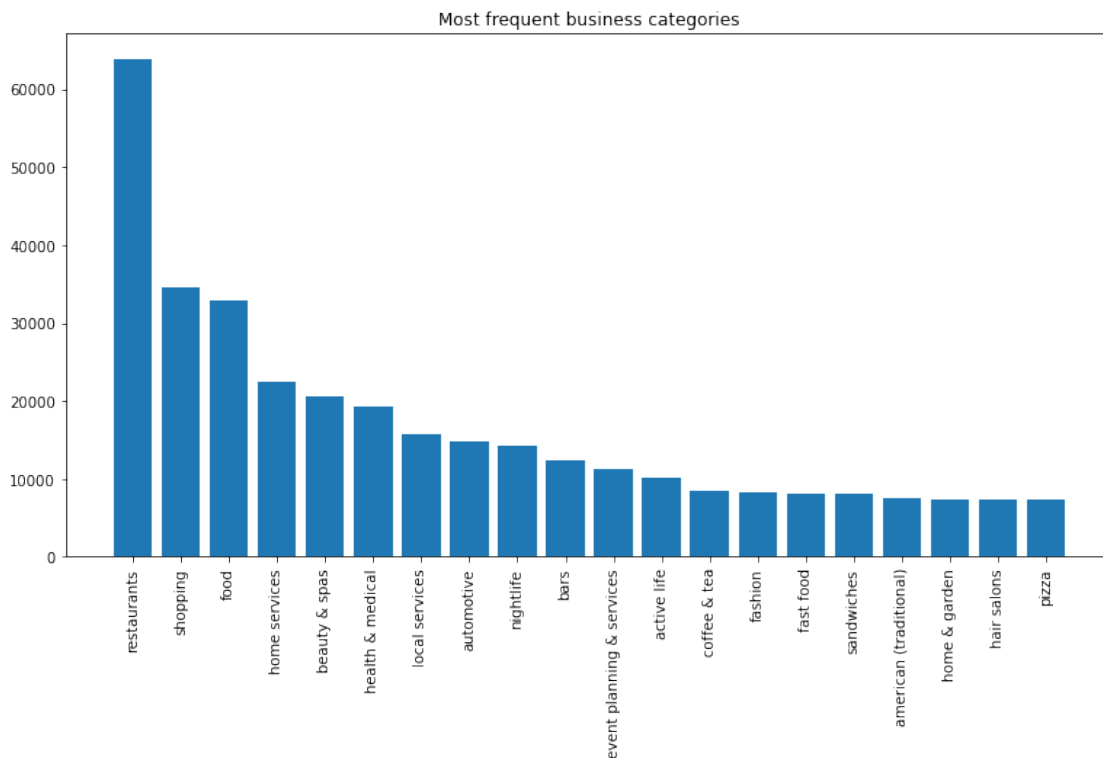
```
[34]: 0      Active Life, Gun/Rifle Ranges, Guns & Ammo, Sh...
1      Health & Medical, Fitness & Instruction, Yoga,...
2      Pets, Pet Services, Pet Groomers
3      Hardware Stores, Home Services, Building Suppl...
4      Home Services, Plumbing, Electricians, Handyma...
...
209388      Japanese, Sushi Bars, Restaurants
209389      Department Stores, Food, Mobile Phones, Fashio...
209390      American (New), Food, Burgers, Restaurants, Fa...
209391      Pet Services, Pet Training, Pets
209392      Tax Services, Professional Services, Accountan...
Name: categories, Length: 209393, dtype: object
```

```
[35]: from collections import Counter
categories = Counter()
df['categories'].str.lower().str.split(', ').apply(categories.update)
print(categories)
```

Let us plot the 20 most frequent business categories.

```
[36]: lists = sorted(categories.items(), key=lambda item: item[1], reverse=True)
x, y = zip(*lists)
```

```
[37]: fig = plt.figure(figsize=(10,5))
ax = fig.add_axes([0,0,1,1])
businesses = x[:20]
frequencies = y[:20]
ax.bar(businesses,frequencies)
plt.xticks(rotation='vertical')
plt.title('Most frequent business categories')
plt.show()
```

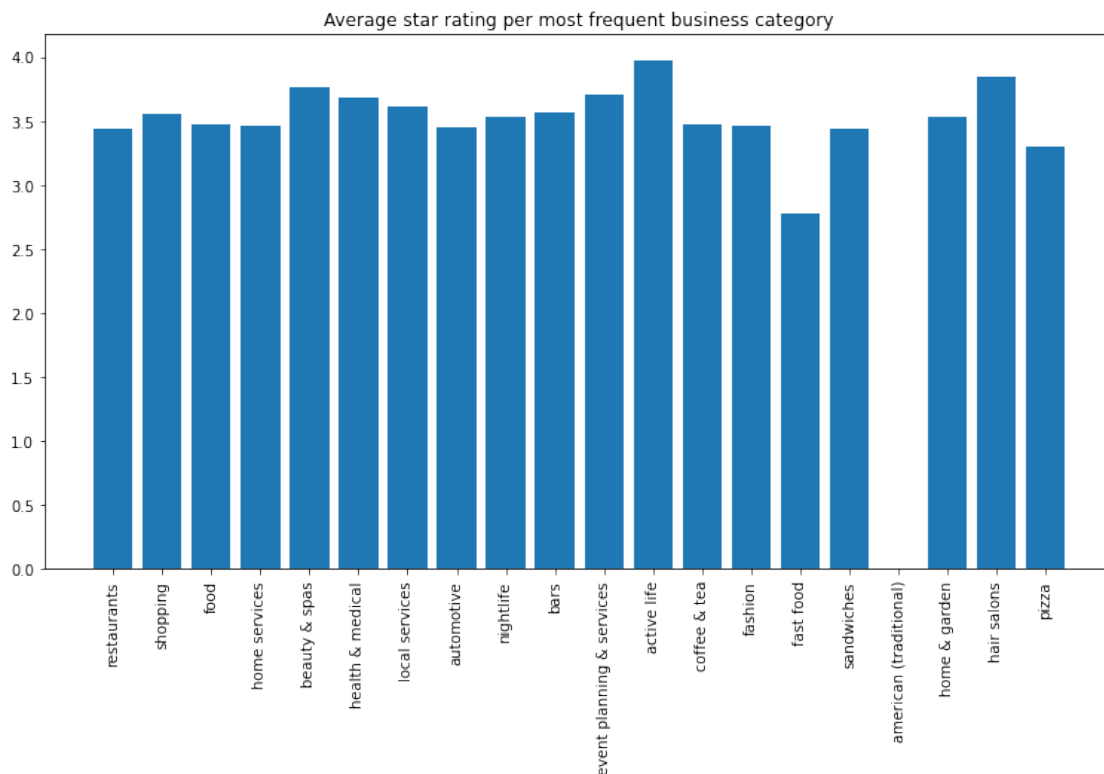


Restaurants, shopping, food and home services among others are the most frequent business categories. Let's see the average star rating these businesses receive.

```
[38]: average_stars = []
      for cat in x[:20]:
          average_stars.append(df[df['categories'].str.lower().str.contains(cat,
→na=False)]['stars'].mean())
```

/opt/conda/lib/python3.7/site-packages/pandas/core/strings/accessor.py:101:  
 UserWarning: This pattern has match groups. To actually get the groups, use  
 str.extract.  
 return func(self, \*args, \*\*kwargs)

```
[39]: fig = plt.figure(figsize=(10,5))
      ax = fig.add_axes([0,0,1,1])
      ax.bar(businesses,average_stars)
      plt.xticks(rotation='vertical')
      plt.title('Average star rating per most frequent business category')
      plt.show()
```

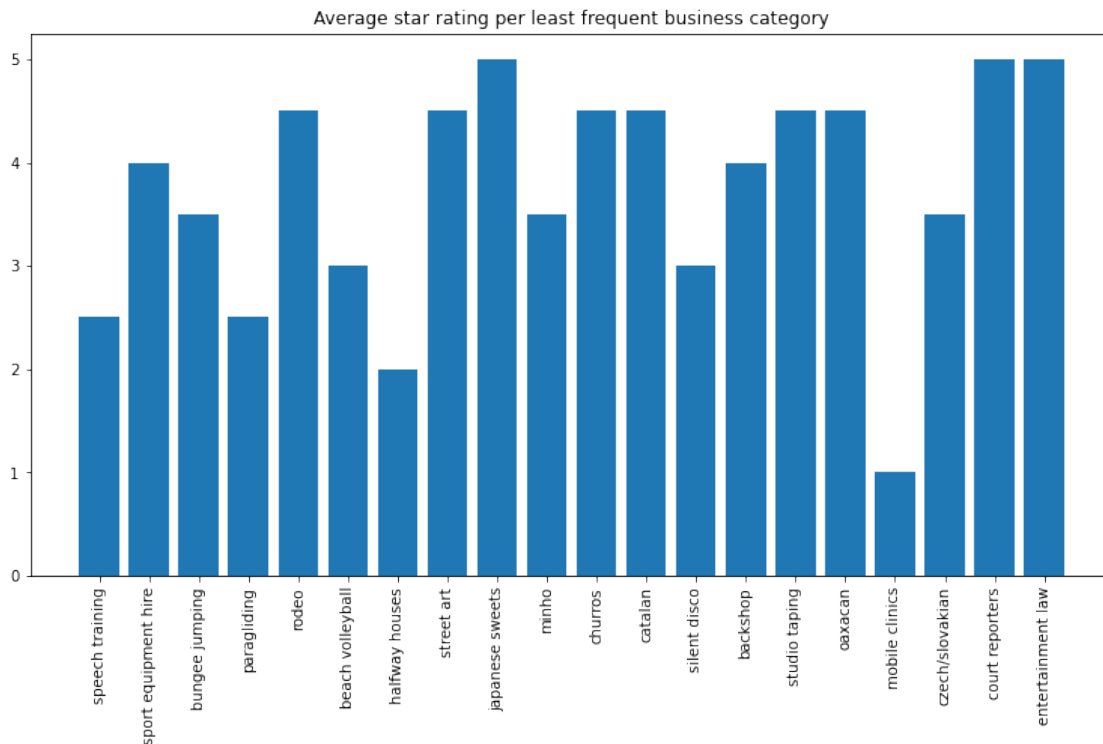


Businesses belonging to categories such as active life, hair salons and beauty and spas are usuall rated higher in average. Let's see the average star ratings for least frequent categories.

```
[40]: average_stars = []
      for cat in x[-20:]:
```

```
average_stars.append(df[df['categories'].str.lower().str.contains(cat,↵
→na=False)]['stars'].mean())
```

```
[41]: fig = plt.figure(figsize=(10,5))
ax = fig.add_axes([0,0,1,1])
businesses = x[-20:]
ax.bar(businesses,average_stars)
plt.xticks(rotation='vertical')
plt.title('Average star rating per least frequent business category')
plt.show()
```



It seems like businesses that are rare get better ratings. Japanese sweets, court reporters and entertainment law categories are highly rated.

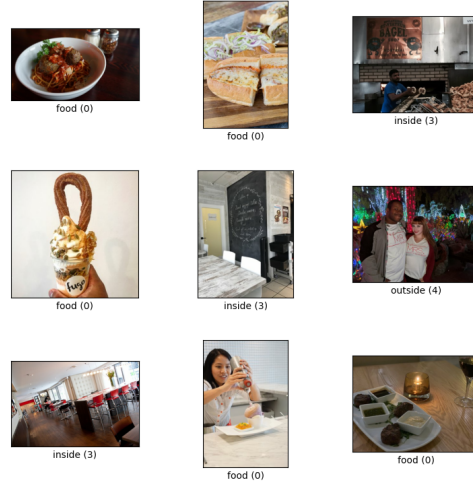
```
[ ]:
```

### 3 Machine Learning Experiments


#### 3.0.1 Image Classification

The first task carried out is image classification on yelp photos dataset. Photos are labeled into five classes including food, inside, outside, drink and menu. An example is shown in Figure 1

A pretrained MobileNetV2 model is used as feature extractor to train a multiclass classification




taking input images and predicting class labels. Details of training could be retrieved from the code accompanying this report. The loss and accuracy plots during training are shown in Figures below.



The trained model is tested on 200 examples from the test set, evaluation results are shown in Figure below.

### 3.0.2 Star Ratings Prediction

The second task carried out is based on reviews and star ratings users have given to businesses. Users write a review text about the business and give a rating ranging from 1 to 5 based on whether they like the place or not. In this section, text processing techniques in Sklearn and a bunch of



	precision	recall	f1-score	support
0.0	0.93	0.99	0.96	126
1.0	0.90	0.98	0.94	45
2.0	1.00	0.58	0.73	19
3.0	1.00	0.38	0.55	8
4.0	0.67	1.00	0.80	2
accuracy			0.93	200
macro avg	0.90	0.78	0.80	200
weighted avg	0.93	0.93	0.92	200

classical machine learning algorithms are used to predict the user's rating based on the review text.

Classifiers such as Random Forest, Nearest Neighbours, AdaBoost and decision trees are evaluated. Results of classifiers on the test set are shown in Figures below.

classifier: Nearest Neighbors				
	precision	recall	f1-score	support
0	0.24	0.12	0.16	2748
1	0.21	0.00	0.01	1637
2	0.13	0.02	0.04	2236
3	0.24	0.25	0.25	4561
4	0.46	0.70	0.56	8818
accuracy			0.39	20000
macro avg	0.26	0.22	0.20	20000
weighted avg	0.32	0.39	0.33	20000



classifier: Random Forest				
	precision	recall	f1-score	support
0	0.64	0.78	0.70	2748
1	0.55	0.05	0.10	1637
2	0.46	0.17	0.25	2236
3	0.44	0.32	0.37	4561
4	0.65	0.90	0.75	8818
accuracy			0.60	20000
macro avg	0.55	0.45	0.43	20000
weighted avg	0.57	0.60	0.55	20000

classifier: AdaBoost				
	precision	recall	f1-score	support
0	0.60	0.67	0.63	2748
1	0.42	0.14	0.21	1637
2	0.41	0.23	0.30	2236
3	0.43	0.35	0.39	4561
4	0.65	0.84	0.73	8818
accuracy			0.58	20000
macro avg	0.50	0.45	0.45	20000
weighted avg	0.55	0.58	0.55	20000

classifier: Decision Tree				
	precision	recall	f1-score	support
0	0.50	0.62	0.55	2748
1	0.17	0.11	0.14	1637
2	0.26	0.22	0.23	2236
3	0.34	0.31	0.32	4561
4	0.64	0.69	0.67	8818
accuracy			0.49	20000
macro avg	0.38	0.39	0.38	20000
weighted avg	0.47	0.49	0.48	20000

### 3.0.3 Recommender Systems

The third class of machine learning experiments conducted are recommendation systems. Based on users and their reviews given to businesses, Memory based recommendation algorithms (Nearest Neighbor) and Model Based Recommendation algorithms (SVD) are used to predict businesses that the user might probably like.

More details of the training and testing are given in the accompanying github [repository](#)